

Programación Científica y Algoritmos Básicos

José R. Dorronsoro
Dpto. de Ingeniería Informática
Escuela Politécnica Superior
Universidad Autónoma de Madrid
28049 Madrid, Spain

Description

- Lecturers: Ana González (coordinator), José Dorronsoro
- Contents:
 - Introduction to Algorithms and Data Structures
 - Dynamic Programming. Levenshtein Distances.
 - Eulerian Paths and DNA Sequencing
 - Algorithm design techniques applied to motif search and sequence alignment
- **First Part:**
 - First three items are a fast review of Chapter 1 in Jones & Pevzner's **An Introduction to Bioinformatics Algorithms**
 - A mix of theoretical classes, lab assignments, problems
 - Evaluation: miniproject + take home

Outline

① Introduction to Algorithms and Data Structures

- Algorithms and Data Structures

- Algorithm Design

- Algorithm Efficiency

② Dynamic Programming

- Revisiting The Change Problem

- DP String Algorithms

③ Euler, Hamilton, and DNA Sequencing

- Eulerian and Hamiltonian Paths

- Hamiltonian and Eulerian DNA Sequencing

Algorithms

① Introduction to Algorithms and Data Structures

- Algorithms and Data Structures

- Algorithm Design

- Algorithm Efficiency

② Dynamic Programming

③ Euler, Hamilton, and DNA Sequencing

Wirth's Equation



- Actually is the title of his book **Algorithms + Data Structures = Programs**
 - Very good book with examples written in Pascal
- Just add input and output to the equation to get useful software
- But first we have to define and clarify what we mean by **Algorithms** and by **Data Structures**

Algorithms



- Many definitions, none too precise
- Adapting from Wikipedia

a set of rules that precisely define a sequence of operations to perform some task and which eventually stop

- Usually written in **pseudocode**:
 - Intermediate between natural language and computer code
 - Not necessarily directly understandable by a computer but close by
 - Simple Python often fine in "simple" tasks
- Three building blocks:
 - **Sequential blocks**
 - **Selections**
 - **Repetitions or Loops**

Sequences, Selections and Loops

- **Sequential blocks:** blocks of (ordinary) sentences that execute sequentially in their entirety



- Sentences may have just straight computations or several calls to functions
- Order of execution according to gravity's law
- In Python: blocks made of sentences with same indentation
- **Selections:** sentences where execution branches to different blocks according to some condition
 - In Python: `if` condition:, `elif` condition:, `else`:
- **Repetitions or loops:** a sentence block is repeated while some condition holds
 - In Python: `while` condition: and also `for` iterations



First Example: Euclid's Algorithm

- Computes the g.c.d. of two positive numbers a, b by repeatedly computing $r = a \% b$ and replacing a by b and b by r while $r > 0$
- In Python:

```
def euclid_gcd(a, b):  
    while b > 0:  
        #print(a, b)  
        r = a % b  
        a = b  
        b = r  
  
    return a
```




- Or more concise (pythonic?):

```
def euclid_gcd(a, b):  
    while b > 0:  
        a, b = b, a % b  
  
    return a
```



But ... Watch Out!!

- In the last version of Euclid's algorithm:
 - What happens if a is 0?
 - What happens if a, b or both are negative?
- Wirth's equation is nice and broadly correct but programs also need quite a bit of **exception handling**, i.e.
 - Detect exceptional situations
 - Tell the program what to do when they arise
- Argument errors are easy to prevent and handle
- Execution exceptions, i.e., things going wrong during execution are harder to detect and prevent
- Programming should be quite **defensive**

Data Structures

- Algorithms work on data
- Single variables are OK for simple algorithms
- **Data structures:** ways to organize more complex data for more advanced algorithms
- Simplest data structures: **strings, lists, arrays**
- More advanced: **dictionaries, sets**
- All of them built in in Python (arrays better through Numpy)
- Advanced structures: **linked lists, trees, graphs**
 - Available in Python through imported modules




Strings, Lists, Arrays and Dicts



- String, list and array elements are accessed through **indices**
 - Similar at first sight (a string is a list of characters?) but in fact quite different
 - Our arrays will be defined and handled by the `numpy` module
- Dicts are made up of `key:value` pairs
- All are Python **objects** with
 - **attributes**: variables with object information and
 - **methods**: functions that act on the object's content
- `dir(object)` lists all of them
 - Try `dir(1)`


Linked Lists

- Made up of individual **nodes** with fields `data`, `next`
 - `data` contains the node's info
 - `next` points to the next node
- They are **dynamic versions** of arrays
- Useful when the number of nodes and/or their location is not known in advance




Trees

- Hold data nodes organized in a **hierarchical** way with a single **root** node and the other ones having a **parent** and perhaps **children**



Graphs

- Made up of **nodes** or **vertices** connected by **edges**
- Possibly the most general data structure: can represent road maps, social networks, protein interactions, ...



Algorithm Design

① Introduction to Algorithms and Data Structures

Algorithms and Data Structures

Algorithm Design

Algorithm Efficiency

② Dynamic Programming

③ Euler, Hamilton, and DNA Sequencing

Writing Algorithms



- Algorithm writing (and programming in general) is usually done on an ad-hoc basis
- It is a **creative act**: must follow programming rules but also requires **imagination, creativity and experience**
 - The same happens with ordinary writing: we cannot fill an empty page just with grammar rules
 - Programming also requires hard work, lots of practice and, also, **quite a bit of algorithm reading**
- Sometimes we can take advantage of general **design techniques**
 - Derived from long experience in problem solving and algorithm analysis
 - Cannot be applied as automatic rules of thumb
 - But may have a wide range of applicability
- We will consider three here: **greedy** algorithms, **divide and conquer** (a.k.a. **recursive**) algorithms and **dynamic programming**

The Change Problem

- Assume we have work as supermarket cashiers and our clients want change in as few coins as possible
- How can we proceed?
- Simplest idea: give at each step the largest coin smaller than the amount that remains to change
- Example: how to give change of 3,44 euros?
 - Easy: one 2 euro coin, one 1 euro coin, two 20 cent coins, two 2 cent coins
- Have to write down the algorithm but the general idea is **greedy**:
 - We try to minimize **globally** the total number of coins
 - Using **locally** at each step the largest coin possible to minimize the amount still to change

The Greedy Change Algorithm

- We will work with an ordered list of coin values, say
1, 2, 5, 10, 20, 50, 100, 200
and save the number of coins of each type used in a dict

```
def change(c):  
    """  
    docstring: write it always!!!  
    """  
    assert c >= 0, "change for positive amounts only"  
  
    l_coin_values = [1, 2, 5, 10, 20, 50, 100, 200]  
    d_change = {}  
  
    for coin in sorted(l_coin_values)[::-1]:  
        d_change[coin] = c // coin  
        c = c % coin  
  
    return d_change
```



- Exercise: pass the coin list as a parameter and add an appropriate assert

Does It Work?



- At first sight yes, but ...
- Try it to give change of 7 maravedís with coin values
1, 3, 4, 5
- What is the answer of the algorithm?
 - Correct answer: just 2 coins, one of 4 maravedís and one of 3
- This often happens with greedy algorithms
 - They are very natural but may give wrong results!!
 - Later we will give a correct algorithm using **Dynamic Programming**, another general algorithm design technique



The Towers of Hanoi

- We are given a set of 64 gold disks of different sizes piled on peg A in increasing sizes, and two other empty pegs B , C
- We want to move the first stack to B one disk at a time using C as an auxiliary peg obeying the rule:
 - **No disk may be placed on top of a smaller disk**
- Easy for 2 disks, not too hard for 3. But for 4 ...???
- Simple **recursive** solution: for N disks
 - Move the first $N - 1$ disks from peg A to C using B as the auxiliary peg
 - Move the remaining disk from A to B
 - Move the $N - 1$ remaining disks from C to B using A as the auxiliary peg
- Very easy to program

Printing Disk Moves

- The following Python code prints the disk moves required:

```
def hanoi(n_disks, a=1, b=2, c=3):  
    assert n_disks > 0, "n_disks at least 1"  
  
    if n_disks == 1:  
        print("move disk from %d to %d" % (a, b))  
    else:  
        hanoi(n_disks - 1, a, c, b)  
        print("move disk from %d to %d" % (a, b))  
        hanoi(n_disks - 1, c, b, a)
```



- But watch out the running times even for small `n_disks`
- In fact the general Hanoi problem is **extremely costly** even for moderate disk numbers

When Will The World End?

- According to a legend, there is large room in a Kashi Vishwanath temple with three posts where 64 golden disks are placed and Brahmin priests keep on moving them
- And the world will end when the last disk move is done
- Q: How many moves will be needed?
- To get an idea
 - Count first the number of moves for small values of N
 - Modify the previous Hanoi code to get a function that returns the number of moves
 - And then decide whether to start worrying
- By the way, the legend is false:
 - The whole thing is a mathematical game devised by Edouard Lucas

Divide And Conquer

- Recursive algorithms often derive from a **Divide and Conquer** strategy:
 - Divide a problem P in M subproblems P_m
 - Solve these separately getting solutions S_m
 - Combine these solutions in a solution S of P
- Two subproblems in Hanoi:
 - P_1 is the subproblem of moving $N - 1$ disks from A to C using B
 - P_2 is the subproblem of moving $N - 1$ disks from C to B using A
 - And we combine the moves according to the Python code
- Efficient algorithms if subproblems are substantially smaller
 - Not the case in Hanoi
- Another clearer example is **binary search** (we'll see it in the exercises)

Algorithm Efficiency

① Introduction to Algorithms and Data Structures

- Algorithms and Data Structures

- Algorithm Design

- Algorithm Efficiency

② Dynamic Programming

③ Euler, Hamilton, and DNA Sequencing

Things To Keep In Mind

- First of all **algorithms must be correct**
 - A fast but wrong algorithm is useless
- It is also desirable that they do not require (much) extra memory
 - The `hanoi` function fulfills this: only its parameters are used
 - Something to watch out in Bioinformatics
- It is also highly desirable that they are as **fast** as possible
 - But ... an algorithm must "read" its inputs
 - If there are many and they are large, the algorithm will likely be slow
 - But desirable execution times should not be far above from the same "order of magnitude" than its inputs' size
- How do measure execution times?



Estimating Execution Times I



- First of all, **forget about just measuring actual times**
 - They depend on the language, the machine, the programmer and, of course, the inputs
 - They are thus too context-dependent to allow meaningful **generalizations**
- We focus instead on **abstract times** measured by counting the **key operations** the algorithm performs on a given input
- For iterative algorithms we usually look for the key operation on the innermost loop
 - Counting how many times these key operations are performed will give us a good estimate of the time their algorithms will take
- This way, the cost of the change algorithm is given by the length of the coin list

Estimating Execution Times II

- The analysis of recursive algorithms is (much) harder
- We can see that possible key operations may be
 - `print("move disk from %d to %d" % (a, b))` for Hanoi
- But while they appear explicitly in the code, they also take place **inside** the recursive calls
- This results in **recurrent estimates** for the cost of recursive algorithms that are often tricky to write and to solve
- Let's try to develop some general strategies on much simpler, loop-based algorithms

Matrix Multiplication

- Well known algorithm: $c_{i,j} = \sum_{k=1}^n a_{i,k} b_{k,j}$ (and quite costly!!)
- Simple (and quite bad) Python code:

```
def matrix_multiplication(m_1, m_2):  
    """ ... """  
    n_rows, n_interm, n_columns = \  
        m_1.shape[0], m_2.shape[0], m_2.shape[1]  
  
    m_product = np.zeros( (n_rows, n_columns) )  
  
    for p in range(n_rows):  
        for q in range(n_columns):  
            for r in range(n_interm):  
                m_product[p, q] += m_1[p, r] * m_2[r, q]  
  
    return m_product
```



- Key operation: clearly $m_1[p, r] * m_2[r, q]$
- How many? Assuming square matrices with N rows/columns, obviously $N \times N \times N = N^3$
 - Substantially larger than problem size $N^2 + N^2 = 2N^2$

Linear Search

- Perhaps the simplest algorithm to search a key in a list:
 - Just compare the key against the list's elements until a match (if any) is found

- Simple Python code:



```
def linear_search(key, l_ints):  
    """ ... """  
    for i, val in enumerate(l_ints):  
        if val == key:  
            return i  
  
    return None
```

- Key operation: clearly `if val == key:`
- Finding `l_ints[0]` requires just one key operation
- Finding `l_ints[-1]` requires $N = \text{len}(l_ints)$ key operations
- If `key` is not in `l_ints`, `linear_search` will also require N key operations



The o , O and Θ Notations I

- It seems that the cost in key operations is given by some function $f(N)$ of input size N
 - We have $f_{MM}(N) = N^3$, $f_{LS}(N) = N$
- We can thus **compare two algorithms** A, B by **comparing their cost functions** f_A, f_B
- We assume that the cost functions are positive and increasing (this should be the case with the abstract execution times of algorithms)
- Given such a pair f, g , we say that $f = o(g)$ if $\frac{f(N)}{g(N)} \rightarrow 0$ when $N \rightarrow \infty$
 - The growth of f is “clearly” smaller than that of g
- Also, $f = O(g)$ if we can find C and N_C s. t. $f(N) \leq Cg(N)$ if $N \geq N_C$
 - g will be bigger than f eventually ($N \geq N_C$) and with help of C

The o , O and Θ Notations II

- Finally, $f = \Theta(g)$ if $f = O(g)$ and $g = O(f)$
- (Very) **Informally** we will understand the preceding as
 - $f < g$ when $f = o(g)$
 - $f \leq g$ when $f = O(g)$
 - $f \simeq g$ when $f = \Theta(g)$ (and, hence, $g = \Theta(f)$)
- O and Θ may be quite messy to check in general, but if $\lim_{N \rightarrow \infty} \frac{f(N)}{g(N)} = L \neq 0$, then $f = \Theta(g)$
- In the preceding examples
 - $N^2 = o(N^3)$ and also, but much less precise, $N^2 = O(N^3)$
 - $\frac{N}{2} = \Theta(N)$ and, also, $N = \Theta\left(\frac{N}{2}\right)$

Complexity of an Algorithm



- From the preceding it seems that given an algorithm A with input I we can “measure” abstract execution times as follows
 - We can identify a **key operation** in A and estimate its abstract work on I by the number $n_A(I)$ of times A executes it on I
 - We can assign a **size** N to the input I
 - We can find a **function** $f_A(N)$ so that $n_A(I) = O(f_A(N))$
- In some cases we will be able to refine this to $n_A(I) = \Theta(f_A(N))$
- We say that $f_A(N)$ gives (is?) the **complexity** of A over entries of size N
- We have $n_{MM}(A, B) = N^3$ with N matrix dimension.
- For Hanoi we have $n_{hanoi}(N) = 2^N - 1$
- For successful searches we have $n_{LS}^e(N) \leq N$, but sometimes $n_{LS}^e(N) = 1$
 - So we say that the **worst** case of LS is $W(N) = N$, as always $n_{LS}(k, I_ints) \leq N$

From Abstract Times to Real Times

- In IPython the magic command `%timeit` allows to estimate execution times of simple functions



```
a = np.ones((100, 100)); b = np.eye(100)
%timeit -n 10 -r 1 matrix_multiplication(a, b)
```

- If it gives us a time estimation of, say, 1 second, what can we expect for matrices with dimensions 500?
 - Since $500 = 5 \times 100$, $500^3 = 125 \times 100^3$, i.e., 125 times bigger
 - Thus, we should expect `%timeit` to report about 125 seconds
- This is not %100 precise, but gives a ball park estimate
- Hence, `matrix_multiplication` is quite costly but `hanoi` is truly awful
- And watch out: straight Python code can be quite slow because of language overheads
 - Heavy duty libraries such as `numpy`, `pandas`, `scipy` or `sklearn` do their work in C or C++ compiled code

Revisiting The Change Problem

- ① Introduction to Algorithms and Data Structures
- ② Dynamic Programming
 - Revisiting The Change Problem
 - DP String Algorithms
- ③ Euler, Hamilton, and DNA Sequencing

Back To Giving Change



- The key trick is to decompose the problem into increasing subproblems and to get a formula to go from one subproblem to the next
- Assume we want to change an amount C with $v_1 = 1, \dots, v_N$ coin denominations
- Let $n(i, c)$ be the minimum number of coins to change an amount c using only the first i coins
 - What we want is $n(N, C)$ which we will get from the easier to understand $n(i, c)$
 - Observe that $n(1, c) = c, n(i, 0) = 0$
- Depending on whether or not coin i enters the change of c we have the following equations

$$\begin{aligned} n(i, c) &= n(i-1, c) && \text{if coin } i \text{ **doesn't enter** the change,} \\ &= 1 + n(i, c - v_i) && \text{if coin } i \text{ **enters** the change} \end{aligned}$$



The DP Change Algorithm

- Therefore, we arrive at

$$n(i, c) = \min\{n(i-1, c), 1 + n(i, c - v_i)\}$$

with a $O(1)$ cost

- And the algorithm is (more or less obviously) correct
- Thus, we simply have to fill the DP matrix but with a cost

$$N \times C \times O(1) = O(N \times C) = O(N \times 2^{\lg C})$$


exponential on the size $\lg C$ of C

- In fact, the change problem is **NP-complete**
 - Nice podcast at **BBC's In Our Time**
- But recall there is a linear complexity algorithm for the Euro coin system

Giving DP Change Algorithm

- Giving change via DP of 7 with coin values 1, 3, 4, 5

cambio de 7 con monedas {1, 3, 4, 5}



DP String Algorithms

- 1 Introduction to Algorithms and Data Structures
- 2 Dynamic Programming
 - Revisiting The Change Problem
 - DP String Algorithms
- 3 Euler, Hamilton, and DNA Sequencing

Editing Strings

- Assume we have two strings and want to edit, i.e., transform one string s_1 into another s_2 according to the following allowable operations
 - Change one character in either string
 - Insert a character in either string
 - Delete a character from either string
- Notice that removing a delete a character from string s_1 is equivalent to adding a character to string s_2
- Also we can keep one of the strings fixed and do all the change/add/delete operations on the other

The Edit Distance

- The **Edit Distance** $d(s_1, s_2)$ between two strings s_1 into another s_2 is the minimum number of edit operations that we have to make to turn, say, s_2 into s_1
- For instance, the edit distance between `unnecessarily` and `unesscessaraly` is 3:

```
unne cessarily  
  a d      c  
un  essessaraly
```



as we just add to the second string an `'n'`, delete an `'s'` and change `'a'` into `'i'`

- We can conceptually substitute the change operation by an insert + delete combination

Approximate String Searching

- These arrangements of character plus blanks are sometimes known as **alignment matrices**
 - By convention, the bottom string is the "target" sequence
- The edit distance can also be used for **approximate string searches**: given a string s find another t in a string list so that their edit distance is minimal
- Below we will penalize these mismatches with a $+1$ scoring penalty
- More generally one can consider scoring matrices δ_{s_i, t_j} associated to particular character mismatches
- Sequence alignment with general scoring penalties are very important in **DNA sequence comparison**
 - Some more details later on and in Jones and Pevzner, Section 6.7



A Dynamic Programming Solution

- Given the full strings S and T with M and N characters respectively, consider the substrings

$$S_i = [s_1, \dots, s_{i-1}, s_i], \quad T_j = [t_1, \dots, t_{j-1}, t_j]$$

- If $d_{i,j}$ is the edit distance between S_i and T_j , we want to find $d_{M,N} = \text{dist}(S, T)$
- Observe that if $s_i = t_j$, then $d_{i,j} = d_{i-1,j-1}$
- And if $s_i \neq t_j$ we have three options
 - Change t_j into s_i ; then $d_{i,j} = 1 + d_{i-1,j-1}$
 - Delete t_j from T_j ; then $d_{i,j} = 1 + d_{i,j-1}$
 - Delete s_i from S_i ; then $d_{i,j} = 1 + d_{i-1,j}$

Filling The DP Matrix

- We thus arrive to the following equations for the Edit Distance problem

$$\begin{aligned}d_{i,j} &= d_{i-1,j-1} && \text{if } s_i = t_j; \\ &= 1 + \min \{d_{i-1,j-1}, d_{i,j-1}, d_{i-1,j}\} && \text{if } s_i \neq t_j\end{aligned}$$

which result in an easy to apply algorithm

- Example: find the edit distance between `biscuit` and `suitcase`
- The cost is clearly $O(M \times N)$, no longer NP-complete but, still, costly
- And an initial memory cost is also $O(M \times N)$, just awful, as problem size is $O(M + N)$
 - But, in fact, quite easy to alleviate if we only care about $d_{M,N}$

An Example

- Example: find the edit distance between `biscuit` and `suitcase`

	\emptyset	b	i	s	c	u	i	t
\emptyset	0	1	2	3	4	5	6	7
s	1	1	2	2	3	4	5	6
u	2	2	2	3	3	3	4	5
i	3	3	2	3	4	4	3	4
c	4	4	3	3	4	5	4	3
e	5	5	4	4	3	4	5	4
a	6	6	5	5	4	4	5	5
s	7	7	6	5	5	5	5	6
e	8	8	7	6	6	6	6	6



Alignment with general penalties

- Sequence alignment with general scoring penalties is very important in DNA sequence comparison
- We can consider different penalties for switching s_i or t_j into the other or for inserting one or the other
 - For switching we may assume a penalty of α_{s_i, t_j}
 - For inserting we may assume a character independent penalty of δ
- We know talk about **alignment costs** $c_{m,n}$ and the new DP penalty equations are

$$c_{i,j} = \min \{ \alpha_{s_i, t_j} + c_{i-1, j-1}, \delta + c_{i, j-1}, \delta + c_{i-1, j} \}$$

where we assume $\alpha_{c,c} = 0$

- This (plus quite a bit of Biochemistry) is at the core of the **Smith–Waterman** and **Needleman–Wunsch** sequence alignment methods
 - More biology-oriented details in Jones and Pevzner, Sections 6.4–6.7

The Longest Common Substring

- Given again strings S, T , we want to find the longest (non necessarily consecutive) common substring (LCS) to both
- As before, let first $\ell_{i,j}$ be the length of the LCS between S_i and T_j . We have now:

$$\begin{aligned}\ell_{i,j} &= 1 + \ell_{i-1,j-1} && \text{if } s_i = t_j; \\ &= \max\{\ell_{i,j-1}, \ell_{i-1,j}\} && \text{if } s_i \neq t_j\end{aligned}$$

- Again this results in an easy to apply algorithm with cost $O(M \times N)$
- Example: find the LCS between `biscuit` and `suitcase`
- Good (and easy) exercise: write down a Python function to find the length of the LCS
- Very good exercise: modify the previous Python function so that it gives one of the possible LCS

Other Problems With DP Solutions




- There is a large number of problems, essentially different and often useful in practice, that can be solved via DP algorithms
- We list some of them:
 - Find the optimal ordering to multiply N matrices
 - Find optimal binary search trees
 - Find RNA's Secondary Structure
 - Cut a rod of a length L in pieces that get as much money as possible when sold
 - Finding the longest palindromic substring of a given string
 - Fit Least Squares segments to data
 - Speech recognition through the Viterbi algorithm
- And many more!

Eulerian Paths

- ① Introduction to Algorithms and Data Structures
- ② Dynamic Programming
- ③ Euler, Hamilton, and DNA Sequencing
 - Eulerian and Hamiltonian Paths
 - Hamiltonian and Eulerian DNA Sequencing

The Bridges of Königsberg

- The bridges of Königsberg (East Prussia) over the Pregel river circa 1700:



- The problem: find a promenade that crosses all bridges but only once
- Tool: a **graph** representation of the problem


Graphs



- Graph: Pair $G = (V, E)$ of a set V of vertices (nodes) and a set E of edges (u, v) with $u, v \in V$
- Edges imply direction: in (u, v) we go from u to v
- In general, graphs are **directed**
- **Undirected** graphs: $(u, v) \in E$ iff $(v, u) \in E$
- **Multigraphs**: graphs with multiple copies of an edge (u, v) or with self-edges (u, u)
- Standard representation: adjacency lists

An Example

- For a graph such as




its adjacency list is

```
v    l_adj
-    -----
1 - 2 - 3
2 - 3
3 - 4
4 - 2
```

- We define $out(w) = \# \text{edges from } w$, $in(w) = \# \text{edges into } w$
- In undirected graphs edges appear “twice”
 - We draw them only once without arrows
 - Then $out(u) = in(u)$ for all $u \in V$ and we just talk of $deg(u)$

The (Multi) Graph of Königsberg

- We can depict the bridges of Königsberg as an **undirected multigraph** (i.e., we allow for multiple bridges/edges between two nodes)



- The problem: find a path that passes **through all edges but only once**
- Such a path in a multigraph is called an **Eulerian path (EP)**

Eulerian Paths on Undirected Graphs

- Let $\pi = \{(u = u_0, u_1), \dots, (u_{K-1}, u_K = v \neq u)\}$ be such an EP
 - When $u = v$ we talk of a circuit
- If $w \neq u, v$ is a node in π , each time we enter w we subtract 1 from $\deg(w)$ and also when we leave w ;
 - Since at the end we have passed through all the edges of w , we must have at the beginning $\deg(w)$ even
- Similarly each time we pass through u inside π we subtract 2 from $\deg(u)$; moreover, we also subtract 1 when we start from u
 - Thus, we must also have $\deg(u)$ odd
- Similarly each time we pass through v we subtract 2 from $\deg(v)$ and we also subtract 1 when we end π at v
 - Thus, we must also have $\deg(v)$ odd
- Thus, **a necessary condition to have an EP is that $\deg(w)$ is even for all w except the first node u and the final one v of π**

Euler's Insight


- Since all the nodes in the previous multigraph have odd degrees, Euler concluded that **no Eulerian path is possible in Königsberg**
- In fact, Euler also proved that the condition is sufficient:
If $\deg(w)$ is even for all nodes w of an undirected graph G , except for two ones u, v , then there is an Eulerian path in G that either starts at u and ends in v or viceversa
- A similar result holds for Eulerian **circuits**:
There is an Eulerian circuit in an undirected graph G if and only if $\deg(w)$ is even for all nodes w

The Bridges of Madrid?

- Madrid has a river with many bridges but no islands
- Does Madrid have a (quite long) Euclidean path? Or perhaps a circuit?
- More generally, think of a city with a river with N bridges and no islands
- Does it have any Eulerian promenade?
- Easy exercise: look at the multigraph and count the degree of its nodes

How to Find an EC

- Assuming an EP exists, the basic idea is start a walk until we cannot go on and re-start the walk if needed
- Example:



- Cost? We simply take out edges, so it is $O(|E|)$

Walking to Find an EC

- The steps are

g) Add L_1 ① all on one even \Rightarrow there is an EC

$$\begin{array}{l} 2 \quad 1 + 2 + 3 \\ 4 \quad 2 + 1 + 3 - 4 - 5 \\ 4 \quad 3 + 1 - 2 - 4 - 5 \\ 4 \quad 4 - 2 - 3 - 5 - 6 \\ 4 \quad 5 - 2 - 3 - 4 - 6 \\ 2 \quad 6 - 4 - 5 \end{array}$$

$\tilde{\pi} = \pi_1: 1^1 - 2^2 - 3^3 - 1^0$

not an EC \Rightarrow we re-start on the residual graph for

② $\pi_2: 2^1 - 4^2 - 3^0 - 5^2 - 2^0$
and we stitch it with $\tilde{\pi} = \pi_1$

$$\tilde{\pi}: 1^0 - 2^0 \quad 2^0 - 3^0 - 1^0$$

$4^2 = 3^0 = 5^2 \Rightarrow$ still not EC
 \Rightarrow re-start from

g) Add L_2

$$\begin{array}{l} 2 \quad 2 + 4 + 5 \\ 2 \quad 3 + 4 + 5 \\ 4 \quad 4 + 2 + 3 - 5 - 6 \\ 4 \quad 5 + 2 - 3 - 4 - 6 \\ 2 \quad 6 - 4 - 5 \end{array}$$

③ $\pi_3: 4^1 - 5^0 - 6^0 - 4^0$ and stitching with

$$\tilde{\pi}: 1 - 2 - 4 \quad 4 - 3 - 5 - 2 - 3 - 1$$

$5 = 6 \Rightarrow$ EC !!!

Hamiltonian Paths



- If G is an undirected connected graph, a **Hamiltonian path** (HP) is a path on G that visits each node **only once**
 - There is a similar definition for Hamiltonian circuits
- Finding HPs may be trivial in some cases, as for complete graphs
- There are also sufficient conditions for special graphs
- But for general graphs, while finding ECs has an $O(|E|)$ cost, finding HPs is much costlier
- In fact, finding HPs in general graphs is another example of an **NP-complete** problem
- Well, and so what??

Hamiltonian and Eulerian DNA Sequencing

- ① Introduction to Algorithms and Data Structures
- ② Dynamic Programming
- ③ Euler, Hamilton, and DNA Sequencing
 - Eulerian and Hamiltonian Paths
 - Hamiltonian and Eulerian DNA Sequencing


DNA Sequencing



- **Note:** this is a very, very light description of one of the many approaches to DNA Sequencing
 - See [DNA sequencing at 40: past, present and future \(Nature, Oct. 2017\)](#) for a very recent review
- Goal: decompose a gene into a sequence of four letters $\{A, C, G, T\}$
- **Shotgun sequencing** broadly follows a four step process:
 - Break the gene into random short “reads” of 100–500 bases
 - Identify read subsequences by hybridizing them on a DNA microarray
 - Reconstruct each read from these subsequences
 - Reconstruct the entire gene from the reads
- First two steps: biochemistry
- Third step: Hamiltonian or (better) Eulerian paths
 - The approach proposed in [An Eulerian path approach to DNA fragment assembly](#) (Pevzner, Tang, Waterman; PNAS, vol. 98, 2001)
- Fourth step: compute the Shortest Superstring Problem (a version of the Traveling Salesman Problem (TSP) , another NP problem)
 - Plus more algorithms and a lot of biochemistry

Shotgun Sequencing

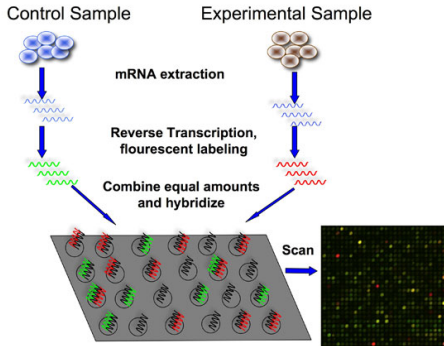
- Idealized hierarchical shotgun sequencing strategy



From Nature

Sequencing by Hybridization

- Scheme of the process:



From bitesizebio.com/7206/introduction-to-dna-microarrays

Microarray Hybridization

- Put all the possible length ℓ probes, i.e., DNA subsequences of a fixed length ℓ , into the spots of a microarray
- Put a drop of fluorescently labeled DNA into each microspot of the array
- The DNA fragment hybridizes with those microspots that are complementary to a certain substring of length ℓ of the fragment
- This way we get all possible length ℓ subsequences that make the fragment but they are **unordered**
 - They follow the order in the microarray but **not the one in the sequence**

ℓ -mers and the Spectrum

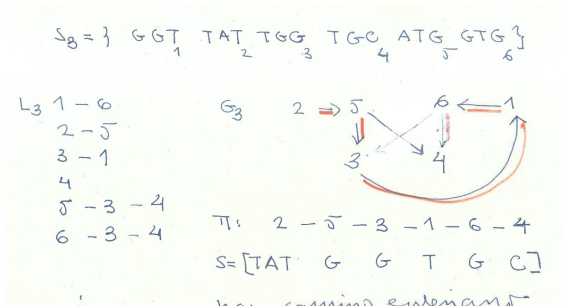
- We call the sequence on each one of the probes an ℓ -mer
- The ℓ -**spectrum** $sp(S, \ell)$ of a sequence S is the set of all the ℓ -mers from S
- For instance, $S = \text{[TATGGTGC]}$ we have
 $sp(S, 3) = \{\text{TAT, ATG, TGG, GGT, GTG, TGC}\}$
- After hybridization, the hybridized probes in the microarray give us **an unordered version of $sp(S, \ell)$** that we have to reorder to recover S
- Define the **overlap** $\omega(s_1, s_2)$ between two ℓ -mers s_1, s_2 as the **longest length of a suffix of s_1 that is also a prefix of s_2**
 - Clearly have $\omega(s_1, s_2) \leq \ell - 1$
- Now, **if s_2 follows s_1 in S , we must have $\omega(s_1, s_2) = \ell - 1$**

Sequencing by Hamiltonian Paths

- We can reconstruct the sequence S by finding an **ordering** s_{i_1}, \dots, s_{i_K} **of** $sp(S, \ell)$ **such that** $\omega(s_{i_j}, s_{i_{j+1}}) = \ell - 1$
- This suggests to define the graph $G_\ell(S) = (V_\ell, E_\ell)$ where
 - $V_\ell = sp(S, \ell)$ and
 - $(s, s') \in E_\ell$ iff $\omega(s, s') = \ell - 1$
- Notice that reconstructing S is equivalent to **pass once through all the nodes of** $G_\ell(S)$
- In other words, **we can reconstruct S by finding a Hamiltonian path in** $G_\ell(S)$

Sequencing by Hamiltonian Paths II

- **Example:** consider $S = [\text{TATGGTGC}]$ and the unordered 3-spectrum $\text{sp}(S, 3) = \{\text{GGT}, \text{TAT}, \text{TGG}, \text{TGC}, \text{ATG}, \text{GTG}\}$
- By inspection, the adjacency list and graph, the HC and the recovered sequence are



Sequencing by Eulerian Paths

- The obvious problem of HP sequencing is the lack of efficient algorithms to solve the HP problem
- The alternative is to try to **have ℓ -mers on the edges** instead of on nodes
- If $s \in sp(S, \ell)$ and s_1 is its $\ell - 1$ prefix and s_2 its $\ell - 1$ suffix, we can consider s as the edge connecting nodes s_1 and s_2
 - Now we have $\omega(s_1, s_2) = \ell - 2$
- We define now the graph $G_{\ell-1} = (V_{\ell-1}, E_{\ell-1})$ where
 - $V_{\ell-1} = sp(S, \ell - 1)$ and
 - $(s, s') \in E_{\ell-1}$ iff they are respectively **prefix and suffix of an** $s \in sp(S, \ell)$
- Notice that now **reconstructing S is equivalent to pass once over all the edges of $G_{\ell-1}$**
- In other words, **we can reconstruct S by finding a Eulerian path in $G_{\ell-1}$**

Eulerian Paths on Directed Graphs

- However, notice that $G_{\ell-1}$ is a **directed** graph and we have to adapt the Eulerian path theory to these graphs
- In an directed graph $G(V, E)$ we have to distinguish between incident and adjacent edges
- For any $u \in V$, we say that (u, v) is an **adjacent** edge for u and an **incident** edge for v
- Recall that the **indegree** $in(u)$ of u is the number of incoming edges to u
- And the **outdegree** $out(u)$ is the number of outgoing edges from u

Eulerian Paths on Directed Graphs II

- Let $\pi = \{(u = u_0, u_1), \dots, (u_{K-1}, u_K = v)\}$, $v \neq u$, be an Eulerian path on G
- If $w \neq u$ is a node in π , each time we enter w we subtract 1 from $in(w)$ and also from $out(w)$ when we leave w
 - Since at the end we have passed through all the edges of w , **we must have at the beginning** $in(w) = out(w)$
- Similarly each time we enter u inside π we subtract 1 from $in(u)$ and also from $out(u)$ when we leave it
 - Moreover, when we start we subtract 1 from $out(u)$
 - We must thus have $out(u) = 1 + in(u)$ **at the starting node**
- Similarly, we must have $in(v) = 1 + out(v)$ **at the ending node**
Thus, we must also have $in(u) = out(u)$
- The above are in fact, **necessary and sufficient conditions to have an EP in a directed G**
- For circuits, i.e., $v = u$, these conditions are $in(w) = out(w)$ for all nodes

Eulerian Sequencing

- Essentially the same $O(|E|)$ algorithm we saw for undirected graphs can be applied to directed ones
- Thus we can efficiently sequence genomic reads
- Example:** consider again $S = [\text{TATGGTGC}]$ and $\text{sp}(S, 2) = \{\text{TA}, \text{AT}, \text{TG}, \text{GG}, \text{GT}, \text{GC}\}$
- Applying the Euler algorithm we obtain

